

# Exercise 3

Jari Mattila - 35260T  
ELEC-E8125 - Reinforcement Learning

October 10, 2021

## Task 1.1

The training performance plots for each of the tasks (Task 1.1 - fixed and GLIE, Task 1.3 - for both initializations, Task 2 - Lunar Lander).

NumPy file `q_values.npy`, which includes the learned Q-values for Task 1.1 for Cartpole with GLIE, saved when the training has finished (don't attach the values for constant epsilon)

NumPy file `value_func.npy`, which contains the value function for the same conditions as in the previous point

Source files: `qlearning.py`, `q_values.npy`, `value_func.npy`

## Task 1.2

The heatmap from the end of the training (Task 1.2).

Plot the heatmap of the value function in terms of  $x$  and  $\theta$ . For plotting, average the values over  $x$  and  $\theta$ .

## Question 1

What do you think the heatmap would have looked like:

- (a) before the training?
- (b) after a single episode?
- (c) halfway through the training?

Justify why for all the cases. Attaching the plots is not required.




Figure 1: This is a sample figure.

## Task 1.3

The training performance plots for each of the tasks (Task 1.1 - fixed and GLIE, Task 1.3 - for both initializations, Task 2 - Lunar Lander).

Source files: qlearning.py

## Question 2

Based on the results you observed in Task 1.3, answer the following questions:

### Question 2.1

In which case does the model perform better?

### Question 2.2

Why is this the case, and how does the initialization of Q values affect exploration?

## Task 2

The training performance plots for each of the tasks (Task 1.1 - fixed and GLIE, Task 1.3 - for both initializations, Task 2 - Lunar Lander).

Source files: `qlearning.py`

## Question 3

Is the lander able to learn any useful behaviour? Why/why not?

# 1 Task 1090

If you add a figure, you can refer to it using Figure. 2.

To cite works, put them in the template.bib file and use [?].

```
1 for episode_number in range(train_episodes):
2     reward_sum, timesteps = 0, 0
3     done = False
4     # Reset the environment and observe the initial state
5     observation = env.reset()
6
7     # Loop until the episode is over
8     while not done:
9         # Get action from the agent
10        action, action_probabilities = agent.get_action(
11            observation)
12        previous_observation = observation
13
14        # Perform the action on the environment, get new state and
15        reward
16        observation, reward, done, info = env.step(action)
17
18        # Store action's outcome (so that the agent can improve
19        its policy)
20        agent.store_outcome(previous_observation,
21            action_probabilities, action, reward)
```




Figure 2: This is a sample figure.